

NumPy

Intro to NumPy: Arrays


- NumPy provides a multidimensional array
 - All elements must be the same data type
 - Many different datatypes supported
 - Size is fixed (memory is allocated for the size specified)
- Arithmetic operations work on arrays
- Provides MANY functions that operate on whole arrays
 - These operations are written in a compiled language, and run fast
 - Generally speaking, **you want to avoid loops to get the best performance.**
 - Can sometimes make code unreadable
- Lots of ways to create arrays

Intro to NumPy: Array Operations

- Arithmetic operator (+, −, /, *) work elementwise
 - $A * B$ is not a matrix product, but instead multiplies the corresponding elements in each array together
 - `dot(A, B)` does a dot product
- Universal functions (`sin`, `cos`, `exp`, ...) work elementwise
- New @ operator
 - Accepted for python 3.5, the “@” will be a new operator in python available for overloading. NumPy will implement it as matrix multiplication
 - <http://legacy.python.org/dev/peps/pep-0465/>
 - $A @ B$ will be equivalent to `np.dot(A, B)`
- Array creation and operations examples...

Intro to NumPy: Array Indexing/Slicing

- Biggest source of confusion: selecting a range is best thought of as referring to the “edges” of the array locations
 - Differs from Fortran, IDL



- For the array above:


- $A[2] = 2$
- $A[2:3] = [2]$
- $A[2:4] = [2 \ 3]$

Note: this same behavior applies to Python lists and strings when slicing

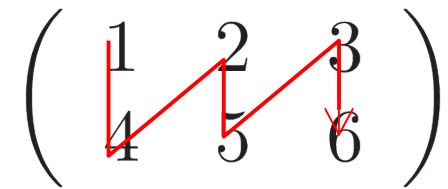
- Note also: **zero-based indexing**

Arrays

- Building block of many numerical methods
- Row vs. Column major: $A(m,n)$
 - First index is called the row
 - Second index is called the column
 - Multi-dimensional arrays are flattened into a one-dimensional sequence for storage
 - Row-major (C, python): rows are stored one after the other
 - Column-major (Fortran, matlab): columns are stored one after the other
- Ordering matters for:
 - Passing arrays between languages
 - Deciding which index to loop over first



Row major



Column major

Arrays

- This is why in Fortran, you want to loop as:

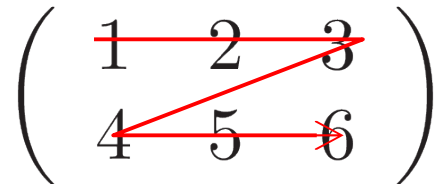
```
double precision :: A(M,N)
do j = 1, N
  do i = 1, M
    A(i,j) = ...
  enddo
enddo
```

- And in C:

```
double A[M][N];
for (i = 0; i < M; i++) {
  for (j = 0; j < N; j++) {
    A[i][j] = ...
  }
}
```

Intro to NumPy: Array Indexing/Slicing

- Remember, multi-dimensional arrays are stored in row-major fashion
 - Rows are stored one after the other, within a row, the column data is closest to one another


$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{pmatrix}$$


- You see this when you print an array:

- `a = numpy.arange(15).reshape(3,5)`

- `print a`

```
[[ 0  1  2  3  4]
 [ 5  6  7  8  9]
 [10 11 12 13 14]]
```

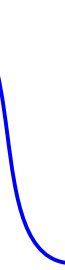


3 rows,
5 columns



- Some slicing examples...

Note that the braces `[]`
show that the columns
are together



Intro to NumPy: Array Views/Copies

- When “copying”, need to understand if two arrays, A and B, point to:
 - the same array (including shape and data/memory space)
 - the same data/memory space (but perhaps different shapes)
 - a separate copy of the data (i.e. stored separately in memory)
- $B = A$ (**assignment**)
 - No copy is made. A and B point to the same data in memory and share the same shape, etc.
- $B = A[:, :]$ (**view** or **shallow copy**)
 - The shape info for A and B are stored independently, but both point to the same memory location for the data
- $B = A.\text{copy}()$ (**deep copy**)
 - The data in B is stored completely separately in memory from A
- **Copying examples...**

Intro to NumPy: Boolean Indexing

- Many fancy ways to index arrays
- $A[A > 4] = \emptyset$
 - Boolean indexing
 - Similar to IDL's where command
- Boolean indexing example...

Avoiding Loops

- Slicing (and using boolean indexing) can be used to avoid loops

More NumPy

- See the tutorial for some other features:
 - Shape manipulation
 - Merging/splitting arrays
 - Fancier indexing
 - Other numpy functions/methods
- [NumPy functions page...](#)